

ACIT4830 – Special Robotics and Control Subject

Topic2 – Classification using Naive Bayes

Evi Zouganeli
OsloMet – Oslo Metropolitan University
(evizou@oslomet.no)

Content

- Probability, joint probability
- Conditional Probability
- Bayes Theorem
- Classification using Naïve Bayes algorithm
- Hands-on exercise


Naive Bayes algorithm

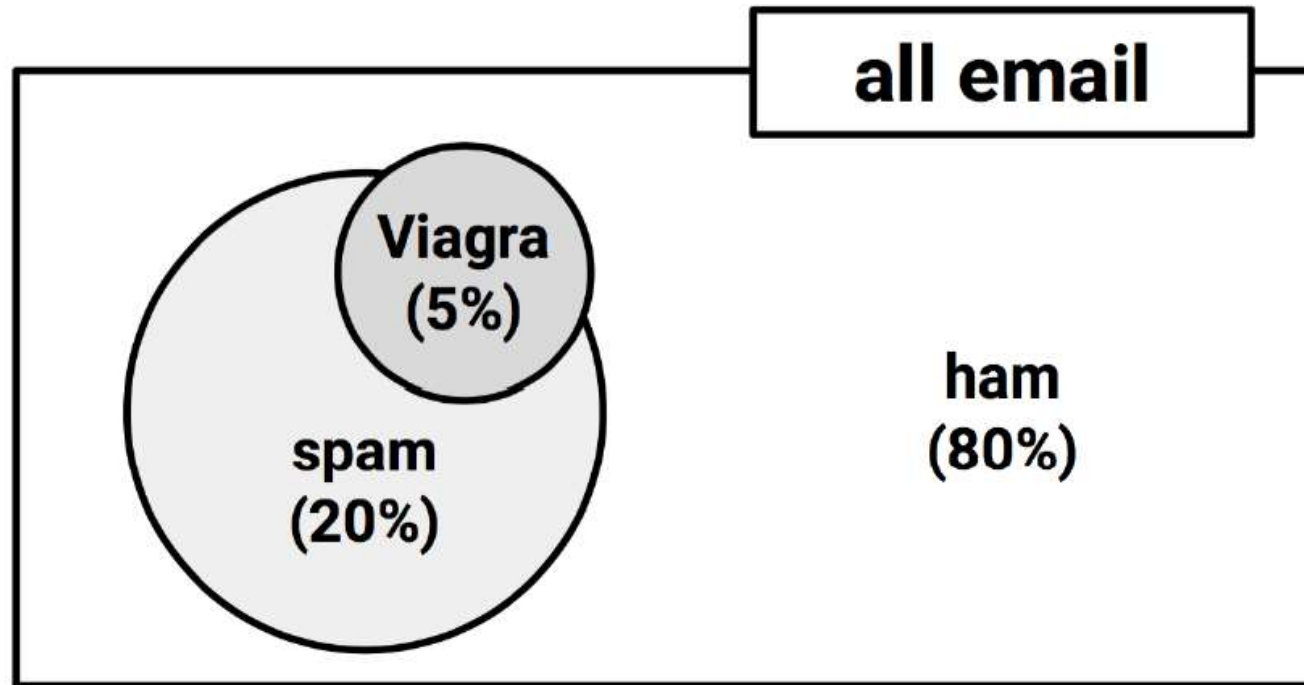
Strengths


- Simple, fast, and very effective
- Does well with noisy and missing data
- Requires relatively few examples for training, but also works well with very large numbers of examples
- Easy to obtain the estimated probability for a prediction

Weaknesses

- Relies on an often-faulty assumption of equally important and independent features
 - Not ideal for datasets with large numbers of numeric features
 - Estimated probabilities are less reliable than the predicted classes
-







Joint Probability A and B

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

Conditional Probability of A given B

A diagram illustrating Bayes' theorem for spam classification. The equation is
$$P(\text{spam} | \text{Viagra}) = \frac{P(\text{Viagra} | \text{spam}) P(\text{spam})}{P(\text{Viagra})}$$
. Four arrows point to different parts of the equation: 'likelihood' points to $P(\text{Viagra} | \text{spam})$, 'prior probability' points to $P(\text{spam})$, 'posterior probability' points to $P(\text{spam} | \text{Viagra})$, and 'marginal likelihood' points to $P(\text{Viagra})$.

likelihood

prior probability

posterior probability

marginal likelihood

$$P(\text{spam} | \text{Viagra}) = \frac{P(\text{Viagra} | \text{spam}) P(\text{spam})}{P(\text{Viagra})}$$

Frequency and Likelihood

	Viagra		
Frequency	Yes	No	Total
spam	4	16	20
ham	1	79	80
Total	5	95	100

	Viagra		
Likelihood	Yes	No	Total
spam	4 / 20	16 / 20	20
ham	1 / 80	79 / 80	80
Total	5 / 100	95 / 100	100

	Viagra (W_1)		Money (W_2)		Groceries (W_3)		Unsubscribe (W_4)		
Likelihood	Yes	No	Yes	No	Yes	No	Yes	No	Total
spam	4 / 20	16 / 20	10 / 20	10 / 20	0 / 20	20 / 20	12 / 20	8 / 20	20
ham	1 / 80	79 / 80	14 / 80	66 / 80	8 / 80	71 / 80	23 / 80	57 / 80	80
Total	5 / 100	95 / 100	24 / 100	76 / 100	8 / 100	91 / 100	35 / 100	65 / 100	100

Naive Bayes

S: Spam

H: Ham (not spam)

B: 'Buy'

C: 'Cheap'

$$P(S | B \cap C) = \frac{P(B|S)P(C|S)P(S)}{P(B|S)P(C|S)P(S) + P(B|H)P(C|H)P(H)}$$

Laplace estimator/ smoothing: correction to ensure that probability terms are non-zero

Word Cloud – R function wordcloud



Data cleaning for text processing - example

- Make all text lower case
- Remove “stopwords” (e.g. and, but), and numbers
- Remove punctuation; replace w/ white spaces if needed
- Keep stem of words
- Remove white spaces

+

- **Document-Term Matrix (DTM)** format: rows are documents, columns are words; and transposed: **TDM**

Data cleaning illustration

SMS messages before cleaning

```
> inspect(sms_corpus[1:3])  
[[1]]  
Hope you are having a good week.  
Just checking in  
[[2]]  
K..give back my thanks.  
[[3]]  
Am also doing in cbe only. But have  
to pay.
```

SMS messages after cleaning

```
> inspect(corpus_clean[1:3])  
[[1]]  
hope good week just checking  
[[2]]  
kgive back thanks  
[[3]]  
also cbe pay
```


Naive Bayes algorithm in R

Naive Bayes classification syntax

using the `naiveBayes()` function in the `e1071` package

Building the classifier:

```
m <- naiveBayes(train, class, laplace = 0)
```

- `train` is a data frame or matrix containing training data
- `class` is a factor vector with the class for each row in the training data
- `laplace` is a number to control the Laplace estimator (by default, 0)

The function will return a naive Bayes model object that can be used to make predictions.

Making predictions:

```
p <- predict(m, test, type = "class")
```

- `m` is a model trained by the `naiveBayes()` function
- `test` is a data frame or matrix containing test data with the same features as the training data used to build the classifier
- `type` is either `"class"` or `"raw"` and specifies whether the predictions should be the most likely class value or the raw predicted probabilities

The function will return a vector of predicted class values or raw predicted probabilities depending upon the value of the `type` parameter.

Example:

```
sms_classifier <- naiveBayes(sms_train, sms_type)
sms_predictions <- predict(sms_classifier, sms_test)
```

Total Observations in Table: 1390

predicted \ actual	ham	spam	Row Total
ham	1202 0.996	28 0.153	1230
spam	5 0.004	155 0.847	160
Column Total	1207 0.868	183 0.132	1390

False positive (arrow pointing to 28)

False negative (arrow pointing to 5)

Confusion Matrix